

CS 280
Spring 2022
Programming Assignment 1

Building a Lexical Analyzer

October 6, 2022

Due Date: Sunday, October 23, 2022, 23:59
Total Points: 20

In this programming assignment, you will be building a lexical analyzer for small programming language and a program to test it. This assignment will be followed by two other assignments to build a parser and interpreter to the same language. Although, we are not concerned about the syntax definitions of the language in this assignment, we intend to introduce it ahead of Programming Assignment 2 in order to show the language reserved words, constants, and operators. The syntax definitions of the small programming language are given below using EBNF notations. However, the details of the meanings (i.e. semantics) of the language constructs will be given later on.

1. `Prog ::= PROGRAM IDENT StmtList END PROGRAM`
2. `StmtList ::= Stmt; { Stmt; }`
3. `Stmt ::= DeclStmt | ControlStmt`
4. `DeclStmt ::= (INT | FLOAT | BOOL) VarList`
5. `VarList ::= Var { ,Var }`
6. `ControlStmt ::= AssigStmt | IfStmt | PrintStmt`
7. `PrintStmt ::= PRINT (ExprList)`
8. `IfStmt ::= IF (Expr) THEN StmtList { ELSE StmtList } END IF`
9. `AssignStmt ::= Var = Expr`
10. `Var ::= IDENT`
11. `ExprList ::= Expr { , Expr }`
12. `Expr ::= LogOExpr ::= LogANDEExpr { || LogANDRexpr }`
13. `LogANDEExpr ::= EqualExpr { && EqualExpr }`
14. `EqualExpr ::= RelExpr [== RelExpr]`
15. `RelExpr ::= AddExpr [(< | >) AddExpr]`
16. `AddExpr ::= MultExpr { (+ | -) MultExpr }`
17. `MultExpr ::= UnaryExpr { (* | /) UnaryExpr }`
18. `UnaryExpr ::= (- | + | !) PrimaryExpr | PrimaryExpr`
19. `PrimaryExpr ::= IDENT | ICONST | RCONST | SCONST | BCONST | (Expr)`

Based on the language definitions, the lexical rules of the language and the assigned tokens to the terminals are as follows:

1. The language has identifiers, referred to by *IDENT* terminal, which are defined as a word that starts by a letter or an underscore ‘_’, and followed by zero or more letters, digits, underscores ‘_’, or ‘@’ characters. It is defined as:

```
IDENT := [Letter _] {(Letter | Digit | _ | @)}  
Letter := [a-z A-Z]  
Digit := [0-9]
```

2. Integer constant is referred to by *ICONST* terminal, which is defined as one or more digits. It is defined as:

```
ICONST := [0-9] +
```

3. Real constant is referred to by *RCONST* terminal, which is defined as zero or more digits followed by a decimal point (dot) and one or more digits. It is defined as:

```
RCONST := ([0-9] +) \. ([0-9] +)
```

For example, real number constants such as 12.0 and 0.2 are accepted, but 2., .2 and 2.45.2 are not.

4. String literals is referred to by *SCONST* terminal, which is defined as a sequence of characters delimited by double quotes, that should all appear on the same line. For example, “Hello to CS 280.” is a string literal. While, ‘Hello to CS 280.’ Or ‘Hello to CS 280.” are not.

5. Boolean constant is referred to by the *BCONST* terminal, which is defined by either the constant value “true” or “false”. Note that the Boolean constants are treated as keywords.

```
BCONST := (true | false )
```

6. The reserved words of the language are: *program, end, print, if, int, float, bool, else, true, false*. These reserved words have the following tokens, respectively: PROGRAM, END, PRINT, IF, INT, FLOAT, BOOL, THEN, ELSE, TRUE, FALSE.

7. The operators of the language are: +, -, *, /, =, (,), ==, >, <, &&, ||, and !. These operators are for plus, subtract, multiply, divide, assignment, left parenthesis, right parenthesis, equality, greater than, less than, logical AND, logical OR, and NOT operations, respectively. They have the following tokens, respectively: PLUS, MINUS, MULT, DIV, ASSOP, LPAREN, RPAREN, EQUAL, GTHAN, LTHAN, AND, OR, and NOT.

8. The semicolon and comma characters are terminals with the following tokens: SEMICOL and COMMA.

9. A comment is defined by all the characters following the sequence of characters “/*” as starting delimiters to the closing delimiters “*/”. Comments may overlap one line, as multi-line comments. A recognized comment is skipped and does not have a token.
10. White spaces are skipped. However, white spaces between tokens are used to improve readability and can be used as a one way to delimit tokens.
11. An error will be denoted by the ERR token.
12. End of file will be denoted by the DONE token.

Lexical Analyzer Requirements:

A header file, `lex.h`, is provided for you. It contains the definitions of the `LexItem` class, and an enumerated type of token symbols, called `Token`, and the definitions of three functions to be implemented. These are:

```
extern ostream& operator<<(ostream& out, const LexItem& tok);  
extern LexItem id_or_kw(const string& lexeme, int linenum);  
extern LexItem getNextToken(istream& in, int& linenum);
```

You MUST use the header file that is provided. You may NOT change it.

- I. You will write the lexical analyzer function, called `getNextToken`, in the file “`lex.cpp`”. The `getNextToken` function must have the following signature:

```
LexItem getNextToken (istream& in, int& linenumber);
```

The first argument to `getNextToken` is a reference to an `istream` object that the function should read from. The second argument to `getNextToken` is a reference to an integer that contains the current line number. `getNextToken` should update this integer every time it reads a newline from the input stream. `getNextToken` returns a `LexItem` object. A `LexItem` is a class that contains a token, a string for the lexeme, and the line number as data members.

Note that the `getNextToken` function performs the following:

1. Any error detected by the lexical analyzer should result in a `LexItem` object to be returned with the ERR token, and the lexeme value equal to the string recognized when the error was detected.

2. Note also that both ERR and DONE are unrecoverable. Once the getNextToken function returns a LexItem object for either of these tokens, you shouldn't call getNextToken again.
3. Tokens may be separated by spaces, but in most cases are not required to be. For example, the input characters "3+7" and the input characters "3 + 7" will both result in the sequence of tokens ICONST PLUS ICONST. Similarly, The input characters

'Hello' 'World', and the input characters 'Hello''World'

will both result in the token sequence SCONST SCONST.

- II. You will implement the id_or_kw() function. Id_or_kw function accepts a reference to a string of a lexeme and a line number and returns a LexItem object. It searches for the lexeme in a directory that maps a string value of a keyword to its corresponding Token value, and it returns a LexItem object containing the keyword Token if it is found and not equal to a TRUE or FALSE, or the identifier token IDENT if not. However, if the string value of a keyword corresponds to either the TRUE or FALSE tokens, the function should return a LexItem object containing the BCONST token instead.
- III. You will implement the overloaded function operator<<. The operator<< function accepts a reference to an ostream object and a reference to a LexItem object, and returns a reference to the ostream object. The operator<< function should print out the string value of the Token in the tok object. If the Token is either an IDENT, ICONST, RCONST, SCONST, or BCONST it will print out its token followed by its lexeme between parentheses. See the example in the slides.

Testing Program Requirements:

It is recommended to implement the lexical analyzer in one source file, and the main test program in another source file. The testing program is a main() function that takes several command line flags. The notations for input flags are as follows:

- -v (optional): if present, every token is printed out when it is seen followed by its lexeme between parentheses.
- -iconst (optional): if present, prints out all the unique integer constants in numeric order.
- -rconst (optional): if present, prints out all the unique real constants in numeric order.
- -sconst (optional): if present, prints out all the unique string constants in alphabetical order
- -bconst (optional): if present, prints out all of the Boolean constants in order.
- -ident (optional): if present, prints out all of the unique identifiers in alphabetical order.

- filename argument must be passed to main function. Your program should open the file and read from that filename.

Note, your testing program should apply the following rules:

1. The flag arguments (arguments that begin with a dash) may appear in any order, and may appear multiple times. Only the last appearance of the same flag is considered.
2. There can be at most one file name specified on the command line. If more than one filename is provided, the program should print on a new line the message “ONLY ONE FILE NAME ALLOWED” and it should stop running. If no file name is provided, the program should print on a new line the message “NO SPECIFIED INPUT FILE NAME FOUND”, and should stop running.
3. No other flags are permitted. If an unrecognized flag is present, the program should print on a new line the message “UNRECOGNIZED FLAG {arg}”, where {arg} is whatever flag was given, and it should stop running.
4. If the program cannot open a filename that is given, the program should print on a new line the message “CANNOT OPEN THE FILE {arg}”, where {arg} is the filename given, and it should stop running.
5. If getNextToken function returns ERR, the program should print “Error in line N ({lexeme})”, where N is the line number of the token in the input file and lexeme is its corresponding lexeme, and then it should stop running. For example, a file that contains an invalid real constant, as 15., in line 1 of the file, the program should print the message:

```
Error in line 1 (15.)
```

6. The program should repeatedly call getNextToken until it returns DONE or ERR. If it returns DONE, the program prints summary information, then handles the flags in the following order:
-sconst, -iconst, -rconst, -bconst and -idents.
The summary information are as follows:

```
Lines: L
Tokens: N
```

Where L is the number of input lines and N is the number of tokens (not counting DONE). If L is zero, no further lines are printed.

7. If the -v option is present, the program should print each token as it is read and recognized, one token per line. The output format for the token is the token name in all capital letters (for example, the token LPAREN should be printed out as the string LPAREN. In the case of the

tokens IDENT, ICONST, RCONST and SCONST, the token name should be followed by a space and the lexeme in parentheses. For example, if the identifier “circle” and a string literal “The center of the circle through these points is” are recognized, the -v output for them would be:

```
IDENT (circle)
SCONST (The center of the circle through these points is)
```

8. The -sconst option should cause the program to print the label STRINGS: on a line by itself, followed by every unique string constant found, one string per line without double quotes, in alphabetical order. If there are no SCONSTs in the input, then nothing is printed.
9. The -iconsts option should cause the program to print the label INTEGERS: on a line by itself, followed by every unique integer constant found, one integer per line, in numeric order. If there are no ICONSTs in the input, then nothing is printed.
10. The -rconst option should cause the program to print the label REALS: on a line by itself, followed by every unique real constant found, one real number per line, in numeric order. If there are no RCONSTs in the input, then nothing is printed.
11. The -bconst option should cause the program to print the label BOOLEANS on a line by itself, followed by every unique Boolean constant found, one Boolean constant per line, in alphabetical order. If there are no BCONSTs in the input, then nothing is printed.
12. The -ident option should cause the program to print the label IDENTIFIERS: followed by a comma-separated list of every identifier found, in alphabetical order. If there are no IDENTs in the input, then nothing is printed.

Note:

You are provided by a set of 18 test case files associated with Programming Assignment 1. Vocareum automatic grading will be based on these testing files. You may use them to check and test your implementation. These are available in compressed archive “PA1 Test Cases.zip” on Canvas assignment. The testing case of each file is defined in the Grading table below.

Submission Guidelines

- 1.1. Submit all your implementation files for the “lex.cpp” and the testing program through Vocareum. The “lex.h” header file will be propagated to your Work Directory.
- 1.2. **Submissions after the due date are accepted with a fixed penalty of 25%. No submission is accepted after Wednesday 11:59 pm, October 26, 2022.**

Grading Table

Case	Test File	Points
1	Successful compilation	1
2	Cannot Open the File (cantopen)	1
3	Empty File (emptyfile)	1
4	Only one file name allowed (onefileonly)	1
5	No Specified Input File (nofile)	1
6	All Integers (ints)	1
7	Bad argument (badarg)	1
8	All Identifiers (idents)	1
9	Valid strings (validstr)	1
10	Invalid string I (invstr1)	1
11	Invalid string II (invstr2)	1
12	Invalid Real I (realerr1)	1
13	Invalid Real II (realerr2)	1
14	Invalid Real III (realerr3)	1
15	Valid operators (allops)	1
16	Invalid Symbol (invsymbol)	1
17	Comments (comments)	1
18	Invalid comment (errcomment)	1
19	All flags set (allflags)	2
		20